

K12 Insight: Template Recommendation Engine

Kaitong Hu, Yanling Sun,
Yichen Peng, Xingchen Wang
November 28, 2023

Meet Our Team



Yanling Sun



Eason Peng



Lulu Wang



Kaitong Hu

K12 Insight pioneers solutions in educational communication

Research shows the need for a strong partnership between stakeholders and schools for the success of pupil's academic performance and educational reforms

Target Customer Group

Primary customers are
K-12 school districts
across the country

Customer Base Size

More than **500** K-12 school districts actively choose K12 Insight as their go-to-partner for educational solutions

Product Portfolio

Let's Talk Communication Platform
Managed Research and Engage Surveys
Professional Development

Let's Talk

In the 2021-22 school year, **87%** of parents reached out to their school district more than once, with 44% of them reached out at least 6 times

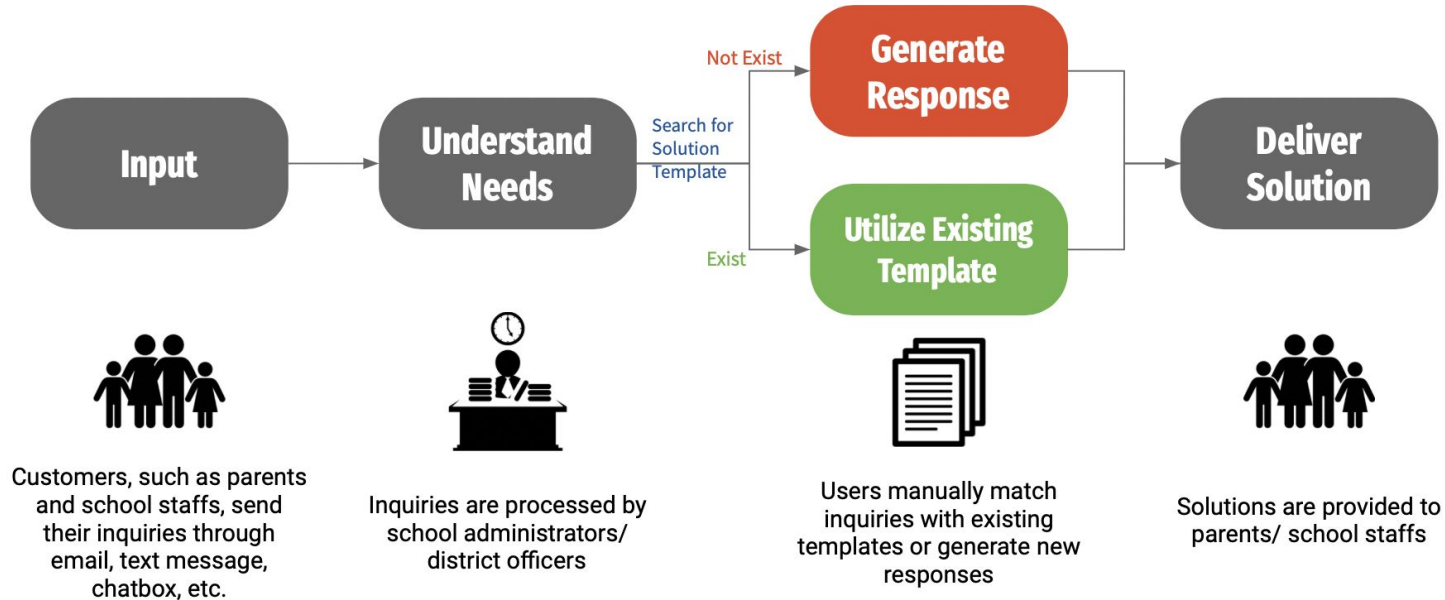


Current customer inquiry handling process needs improvements

■ Time-Consuming Selection Process

■ Overlooked Relevant Templates

■ Inaccurate Keyword-Based Searches



Our recommendation engine leverages NLP & LLMs to automate the template matching process

Sign Up

Sat, Dec 2, 13:30pm (3 Minutes Old) 



ORIGINAL MESSAGE

Hello, how do I sign up for an account so I can access my student's information in the Aspen Family Portal?

ENTER CUSTOMER REPLY

Subject Re: Sign Up 

Cc Response Template

Hello [[Customer's Name]],

Thank you for contacting the YCSD IT department. The first step is to get an **Aspen Family Portal** activation letter. If you have not gotten a portal activation letter, then you must contact your school. Schools issue the initial portal activation letter. If you have the portal activation letter then you simply **sign up for an account** using a computer, laptop or desktop. You cannot sign up for an account using a mobile device. (This is explained in the portal activation letter.) Once you have created an account, you can use either a computer or mobile device to access your student's information in Aspen.

This project aims to transform educational communication

Enhanced Communication

Effective communications is crucial for building positive relationships between schools and students

Workflow Automation

Help streamline and automate processes for managing messages, which saves time and ensures that messages are handled consistently



Trust-Building

Trust is built through transparent and effective communication, responsiveness, and engagement

Stakeholder Engagement

Timely and appropriate responses help fosters a sense of collaboration, leading to improved stakeholder engagement and support for school initiatives

PROJECT WORKFLOW

Data Cleaning & EDA

Perform initial preprocessing of the data to ensure it is clean and formatted correctly. Conduct EDA to understand the distribution of key variables

Text Encoding

Use a Sentence Transformer model to encode sentences into high-dimensional vectors while capturing the semantic meaning of sentences

Similarity Matching

Employ various similarity indexes to measure the similarity between different text vectors

Model Building

Test multiple machine learning models and perform hyperparameter tunings to find the most effective model configuration

Evaluation

Evaluate the model using appropriate metrics

Three main tables are joined for modeling

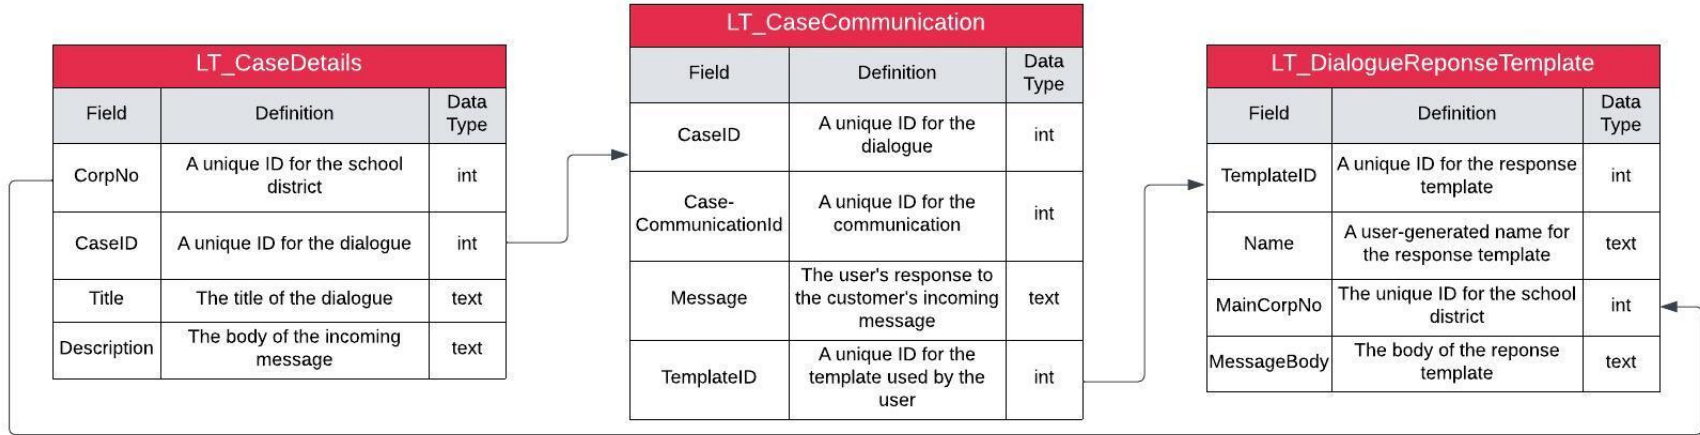
Key Terms:

Customer: The person sending an incoming message to the Let's Talk application

User: An employee of the school district who is responding to the customer's message

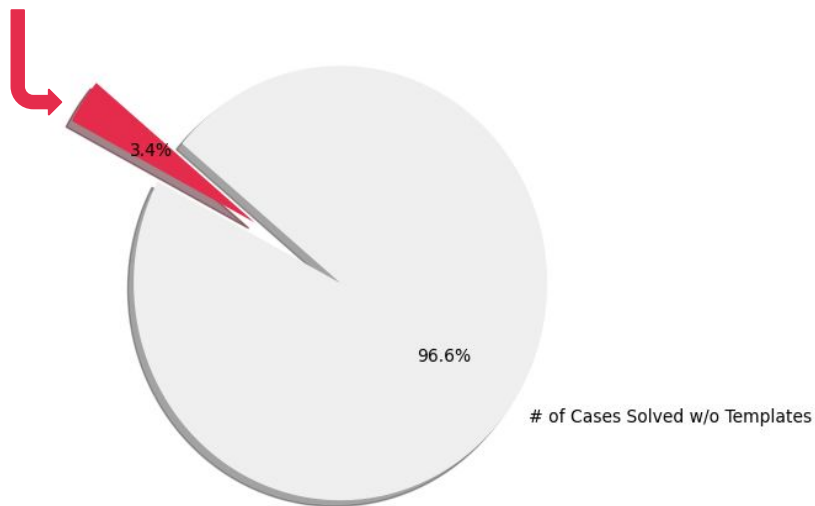
Dialogue: A dialogue represents a single interaction between a customer and the user(s)

Communication: A user's response to the customer




The template feature is widely overlooked by platform users

Among all cases, only a small fraction of them are resolved using pre-written response templates



Even school districts that receive the most number of cases underutilize text templates



Among 1.7M cases, only 57K of them can be used for model training

Implementing a hybrid approach by computing multiple scores for each query-template pair

Major Components

Collaborative Filtering

Enhancing the relevance of template matching by integrating **insights from similar past queries and their corresponding template usage**

	1	2	3	4	5	6
a	5		1	1		2
b		2		4		4
c	4	5		1	1	2


Content-Based Filtering

The **dot product** between the embeddings of customer inquiries and text templates is used to quantify the degree of similarity



Text Similarity

Jaccard Similarity score is used to evaluate how closely an inquiry aligns with a predefined text template



Fine-tuning Sentence Transformer proofs to generate more precise and contextually relevant embeddings

Multi-qa-mpnet-base-dot-v1 is fine-tuned to address the complex nature of customer inquiries

Inquiry


Request to join **Classroom Teachers Group**.

Response

To prevent students from incidentally signing up to Google Classroom as a teacher, and teachers incidentally signing up as Students, Google Classroom membership is a GIPS managed Group. You recently applied for membership as a **Google Classroom Teacher**. Your Teacher application has been reviewed and accepted. When you next log into Google Classroom, you may enter as a GIPS Teacher.

Similarity Score

Before fine-tuning
24.15

Now
26.57 

All scoring components are used jointly for downstream classification tasks

Scoring components for Query M

Collaborative Filtering Result

Template 4535	Template 1481	Template 1456
0.0000	0.5589	3.8836

Content-Based Result

Template 4535	Template 1481	Template 1456
10.5696	13.7951	13.1870

Text Similarity Result

Template 4535	Template 1481	Template 1456
0.0232	0.0540	0.0361

Similarity Score	Collaborative Score	Jaccard Score	Matched
10.5696	0.0000	0.0232	0
13.7951	0.5589	0.0540	0
13.1870	3.8836	0.0361	1
...			



Logistic Regression



Neural Network


dmlc
XGBoost

A combination of fine-tuned embeddings & XGBoost demonstrates the best classification result

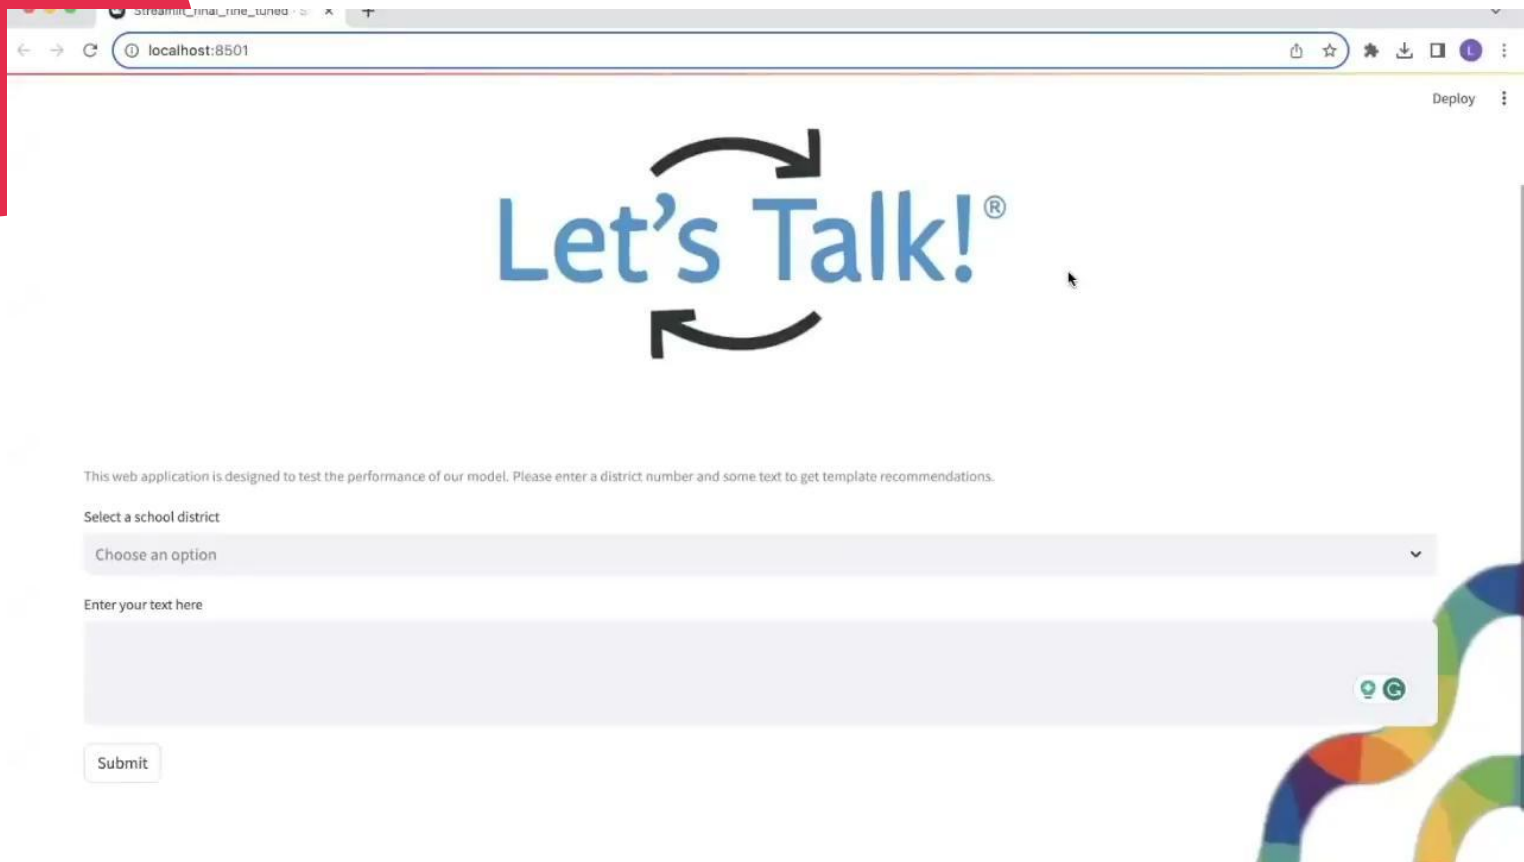
Model Performance on District 45671

Recall Rate		
	No Fine-tuning	Fine-tuning
XGBoost	0.71	0.80
Logistic Regression (LR)	0.64	0.69
Neural Network (NN)	0.64	0.60

Number of Passed Test Cases Across Districts (Out of 20)



Streamlit demonstration



Streamlit_demo_tuned

localhost:8501

Deploy

Let's Talk!®

This web application is designed to test the performance of our model. Please enter a district number and some text to get template recommendations.


Select a school district

Choose an option

Enter your text here

Submit

Our solution delivers more than just efficiency



Template Relevancy

90% relevant templates are delivered to 70% of school districts, reducing response time by 60%

Operational Costs

Reduce annual operational costs by **\$0.7M**

Customer Base

Enhanced efficiency in addressing inquiries attracts more school districts, targeting a **60%** expansion in customer base

Future work to elevate model performance



Collect feedbacks from users, including explicit ratings, comments as well as implicit feedback drawn from user interactions and behaviors within the system.



Continuously update the model by adding new scoring components and try alternative classification algorithms for better performance



Develop a model to create templates for frequently used answers that haven't yet been templated.

Our sincere thanks and appreciation to



**for the opportunity to contribute to the
field of educational communication**

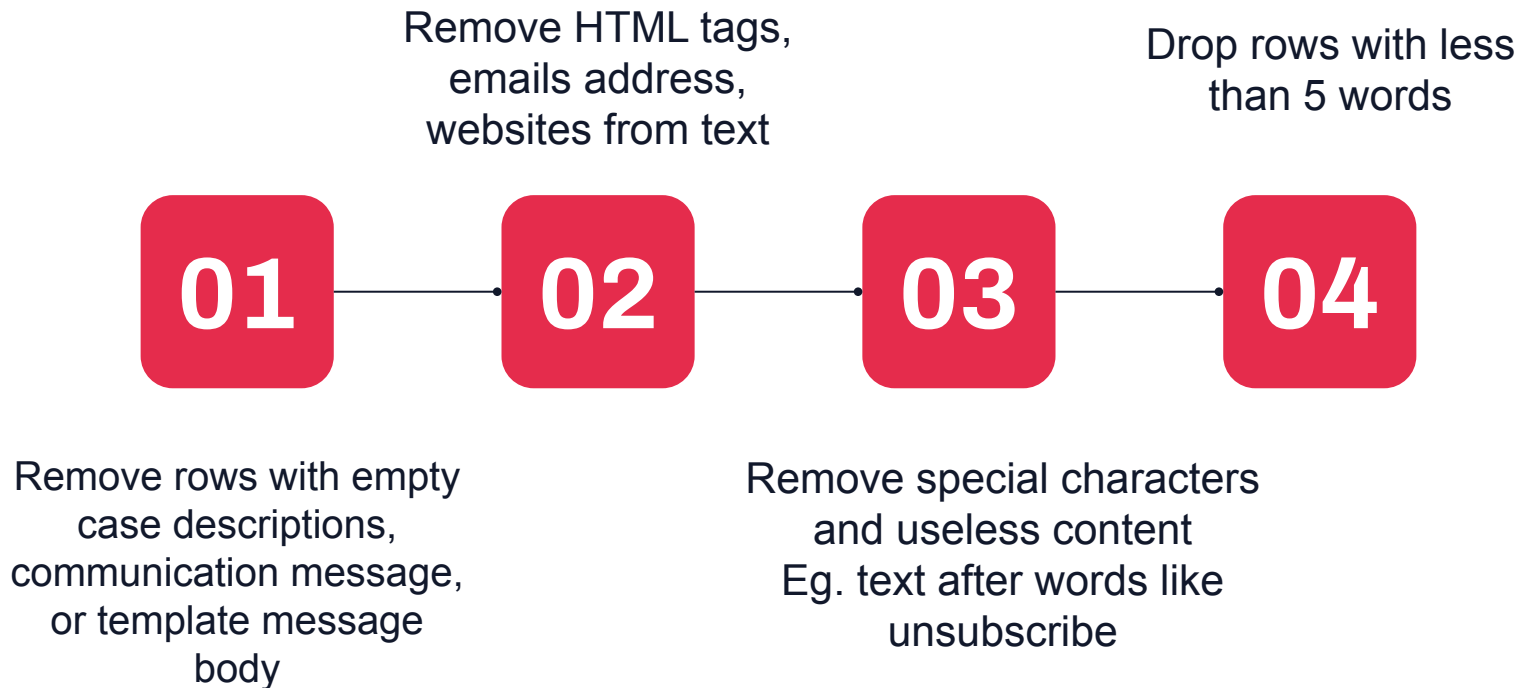




Any Questions?

Appendix

DATA CLEANING



MODELING - SCORE COMPONENTS

Score = queryM-templateN similarity

+ \sum (top10 similar queries of queryM * interaction with templateN)

+ Text Similarity

+ ...

QueryM-TemplateN Similarity

Definition:

- The dot product of query and template

Models:

- Multi-QA-MPNet-Base-Dot
- Multi-QA-DistilBert-Cos
- Multi-QA-MiniLM-L6-Cos

\sum (Top10 Similar Queries of QueryM* Interaction with TemplateN)

Definition:

- Measure the similarity of this template with queries similar to QueryM

Steps:

- Identify top 10 similar queries for QueryM
- For each template, calculate the total # of times it was used among top 10 relevant templates

Text Similarity

Definition:

- Measure how closely an inquiry aligns with a predefined text template

Methods:

- Jaccard Similarity Score

MODELING - SENTENCE TRANSFORMER

Pretrained models from sentence_transformers will be used to encode queries and pre-written templates into high-dimensional word embeddings for similarity matching.

Q&Q Matching with 'all-mpnet-base-v2'

To identify semantically similar customer inquiries within the database. This component aids in understanding the effectiveness of past responses and in optimizing future responses.

- Scoring: Dot-Product

Q&A Matching with 'multi-qa-mpnet-base-dot-v1'

To compare the semantic content of customer inquiries against a set of pre-written text templates. This approach aims to provide direct and accurate responses to customer inquiries.

- Scoring: Dot-Product

MODELS AND ERROR METRICS

Candidate Models


- ▶ Logistic Regression
- ▶ Feedforward Neural networks
- ▶ XGBOOST

Choices of Error Metrics

- ▶ Recall
- ▶ Precision
- ▶ F1 Score

Model Selection Criteria

- ▶ High Recall: True template is recommended
- ▶ Precision Not Too High: Up to five templates are recommended



$$\begin{aligned} \text{F1 Score} &= \frac{2}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}} \\ &= \frac{2 \times \text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \end{aligned}$$

ANNUAL COST REDUCTION CALCULATION

Hours saved per day:

$$\begin{aligned} & (\text{original response time per query}(\text{min}) - \text{new response time per query}(\text{min})) / 60 * \text{number of queries per day} \\ & = (5 - 2) / 60 * 2700 = 135\text{H} \end{aligned}$$

Annual cost reduction:

$$\begin{aligned} & \text{Hours saved per day} * \text{working day per year} * \text{salary per hour} \\ & = 135 * 264 * 20 = \mathbf{712,800 \sim \$0.7M} \end{aligned}$$